

Towards Decoupling Application Descriptions from their Execution Environments

Journées TARANIS, Lyon

Quentin GUILLOTEAU¹, Hélène COULLON², Christian PEREZ¹

2025-10-17

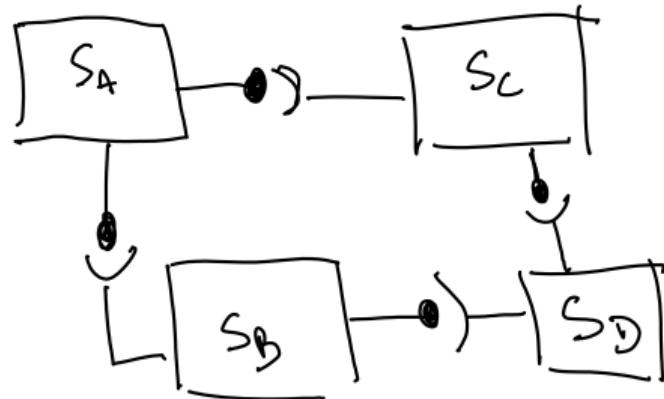
¹Inria

²IMT Atlantique



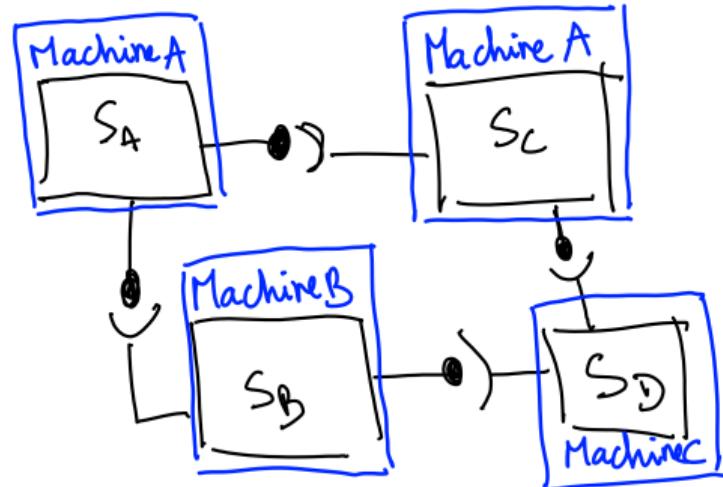
Context – Deployment of Cloud Applications

- deployment of cloud applications
- application = set of (micro)services with interfaces, interconnected
- can be modeled as black box components with ports (e.g., USE/PROVIDE) for better composability
- connected components = assembly



Context – Deployment of Cloud Applications

- deployment of cloud applications
- application = set of (micro)services with interfaces, interconnected
- can be modeled as black box components with ports (e.g., USE/PROVIDE) for better composability
- connected components = assembly



→ Software components need to be deployed onto “**machines**”

Resources, Provisioning, and Config. Management

What even are “machines”?

Does not have to be machines ↪ VMs, containers, baremetal, etc.

↪ “Resources” ↪ Execution units

Usually:

1. Provisioning (e.g., via Terraform):
 - Descriptions of the resources
 - Manual, depends on the tool used
2. Configuration Management (e.g., with Ansible):
 - Installation/configuration/etc.
 - Might depend on the resources

Resources, Provisioning, and Config. Management

What even are “machines”?

Does not have to be machines ↪ VMs, containers, baremetal, etc.

↪ “Resources” ↪ Execution units

Usually:

1. Provisioning (e.g., via Terraform):
 - Descriptions of the resources
 - Manual, depends on the tool used
2. Configuration Management (e.g., with Ansible):
 - Installation/configuration/etc.
 - Might depend on the resources

↪ Mapping between the resources and the components?

Motivations and Research Statement

Research Statement

How to model component assemblies to support their deployment in **various execution contexts?**

Motivations and Research Statement

Research Statement

How to model component assemblies to support their deployment in **various execution contexts?**

Motivations – Portability

- Application developer not responsible to execution context
- Avoiding vendor-locking
- Multi-cloud deployments
- Resource heterogeneity

Motivations and Research Statement

Research Statement

How to model component assemblies to support their deployment in **various execution contexts?**

Motivations – Portability

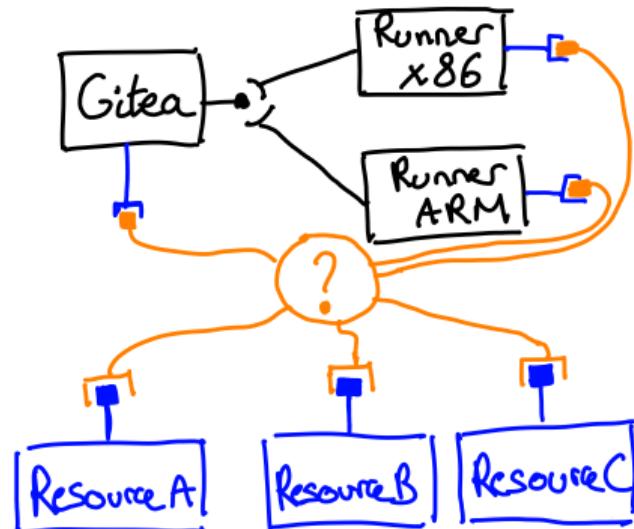
Motivations – Unavailable Services / Urgent Computing

- Resource provider has some issue
- Some resources or services are not available
- but need to deploy the application right now, with whatever resources available

Modelisation

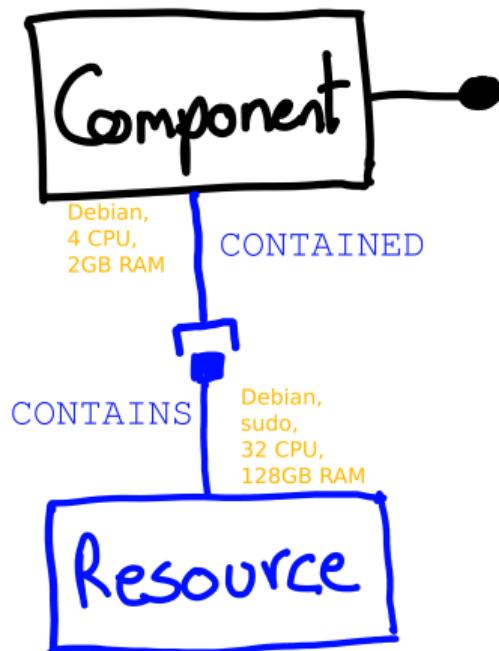
Model – The Main Idea

~~~ Model resources as components



## Model – Components, ResourcePorts, and Attributes

- Components have the classical USE/PROVIDE ports
- but also ResourcePorts
  - port types (eq. to use/provide):  
**CONTAINED/CONTAINS**
  - models where a component is deployed
  - One port = one resource
  - One component can model several resources
- Set of **functional** attributes linked to the ports  
(e.g., CPU, MEMORY, etc.)
- Set of non-functional attributes linked to the ports  
(e.g., FLOPs, CO2/W, etc.)



## Model – Instances and Assemblies

### Instances

Instance = Component Type + Identifier

(For both applicative and resource components!)

Component Type = Ports + ResourcePorts + Attributes

# Model – Instances and Assemblies

## Instances – Component Type + ID

### Assembly

- Set of component instances
- Port Connections
  - between application ports
  - but also between ResourcePorts

## Model – Instances and Assemblies

Instances – Component Type + ID

Assembly – Set of Component Instances + Port Connections

Incomplete Assembly

Assembly where at least one ResourcePort of type **CONTAINED** is not connected

Complete Assembly

Assembly where all the ResourcePorts of type **CONTAINED** are connected

→ A complete assembly is “deployable”

## Model – Instances and Assemblies

Instances – Component Type + ID

Assembly – Set of Component Instances + Port Connections

Incomplete Assembly

Assembly where at least one ResourcePort of type **CONTAINED** is not connected

How to go from an incomplete to a complete assembly?

Complete Assembly

Assembly where all the ResourcePorts of type **CONTAINED** are connected

→ A complete assembly is “deployable”

## Mapper – Definition

### Mapper

$\text{Mapper} : \text{Assembly} \rightarrow \text{Assembly}$

- Only manipulates [ResourcePorts](#)
- Does not have to return a complete assembly (i.e., can solve a sub-problem)

# Mapper – Definition

## Mapper

$\text{Mapper} : \text{Assembly} \times \text{Universe} \rightarrow \text{Assembly}$

- Only manipulates **ResourcePorts**
- Does not have to return a complete assembly (i.e., can solve a sub-problem)

## Universe

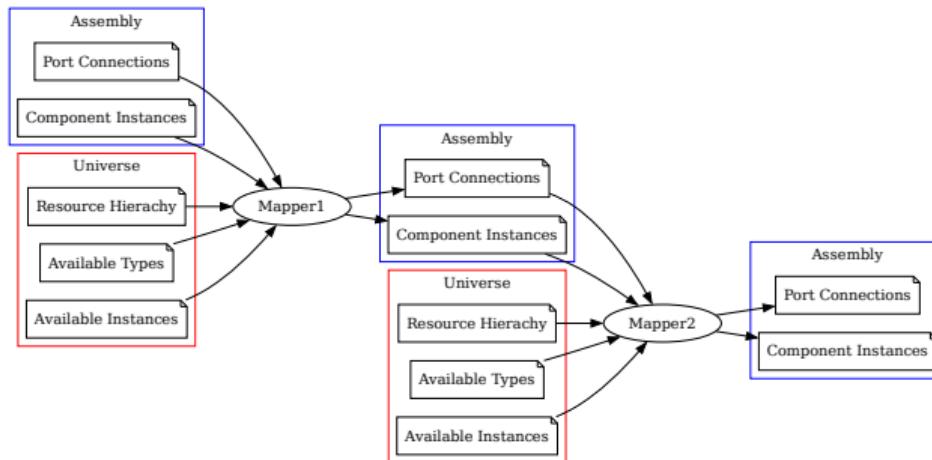
- Set of resource component instances: resources already available
  - Set of resource component types: resources that can be instantiated
  - A relation to model if a resource can be deployed onto another
- Can be generated based on the state of the platform(s)

# Mapper – Definition

## Mapper

$$\text{Mapper} : \text{Assembly} \times \text{Universe} \rightarrow \text{Assembly}$$

Universe – Set of resource instances and types to complete assemblies



## Mapper – Examples of Mappers

**Warning!**: We are **not** proposing new placement algorithms!

### Simple Greedy Mapper

- for each unconnected **CONTAINED** ports
  - go through the set of resources instances and see if one can connect
  - if not, go through the set of resources types, and see if one can connect, if so instantiate this resource
- continue until no **CONTAINED** port is left unconnected (might need to backtrack)

### Integration with Constraint Solvers (via Minizinc)

- Difficult to express our problem as a simple variant of knapsack/bin-packing ...
- A “flat” representation of the problem
- Have to “instantiate” the resource types to feed the solver: how many?

# Mapper – Examples of Mappers

**Warning!**: We are **not** proposing new placement algorithms!

## Simple Greedy Mapper

- for each unconnected **CONTAINED** ports
  - go through the set of resources instances and see if one can connect

Ok, we have a complete assembly, now what ?

- continue until no **CONTAINED** port is left unconnected (might need to backtrack)

## Integration with Constraint Solvers (via Minizinc)

- Difficult to express our problem as a simple variant of knapsack/bin-packing ...
- A “flat” representation of the problem
- Have to “instantiate” the resource types to feed the solver: how many?

## Assembly – Transformation into Concerto

---

We need a way to **execute/deploy** our assemblies

## Assembly – Transformation into Concerto

---

We need a way to **execute/deploy** our assemblies  
→ We choose Concerto

## Assembly – Transformation into Concerto

We need a way to **execute/deploy** our assemblies  
→ We choose Concerto

### Why Concerto?

- Performance!
- and support for reconfiguration (cf. perspectives)

### Transformation into Concerto

- We transform the Assembly in our model into a Concerto assembly
- All of our components have a Concerto implementation in Python
- (behavior support = “only pushB(start)”)

## **Implementation / Evaluation**

---

## Implementation / Evaluation

---

- Realization of our Model in Nickel<sup>1</sup> (JSON + types + contracts +  $\lambda$ )
- Slight changes of Concerto (new port type, add parameters to components)
- Implemented several resource components (i.e., in our model and in Concerto)
  - Grid'5000 (basic env.), Kadeploy,
  - QEMU (x86, arm64), VM on Grid'5000,
  - Docker,
  - K3S cluster (control node + agents)
  - preallocated resource (e.g., laptop)

---

<sup>1</sup><https://nickel-lang.org/>

# Scenario – Gitea and CI Runners

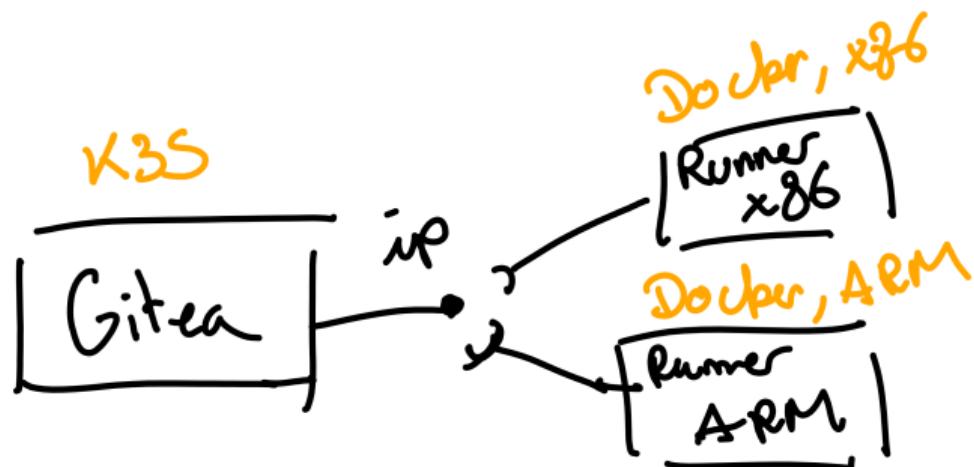
## Assembly

- Gitea: Git hosting service
  - K3S (http, DB, storage)
- CI runners
  - Docker container
  - x86 and arm64

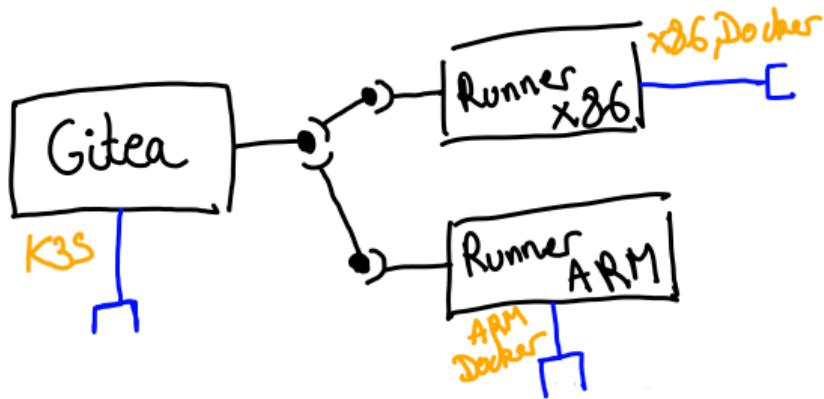
## Scenario

Deploy the same assembly ...

- ... in production (Grid'5000)
- ... in “incomplete” prod.
- ... on my laptop (dev)

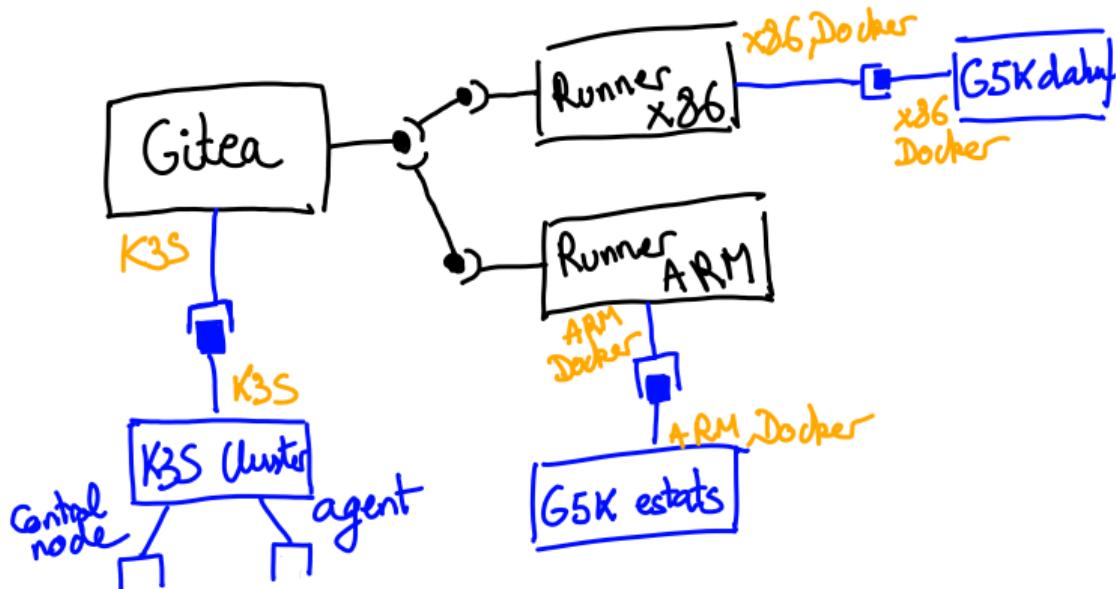


## Scenario – Gitea and CI Runners – Production



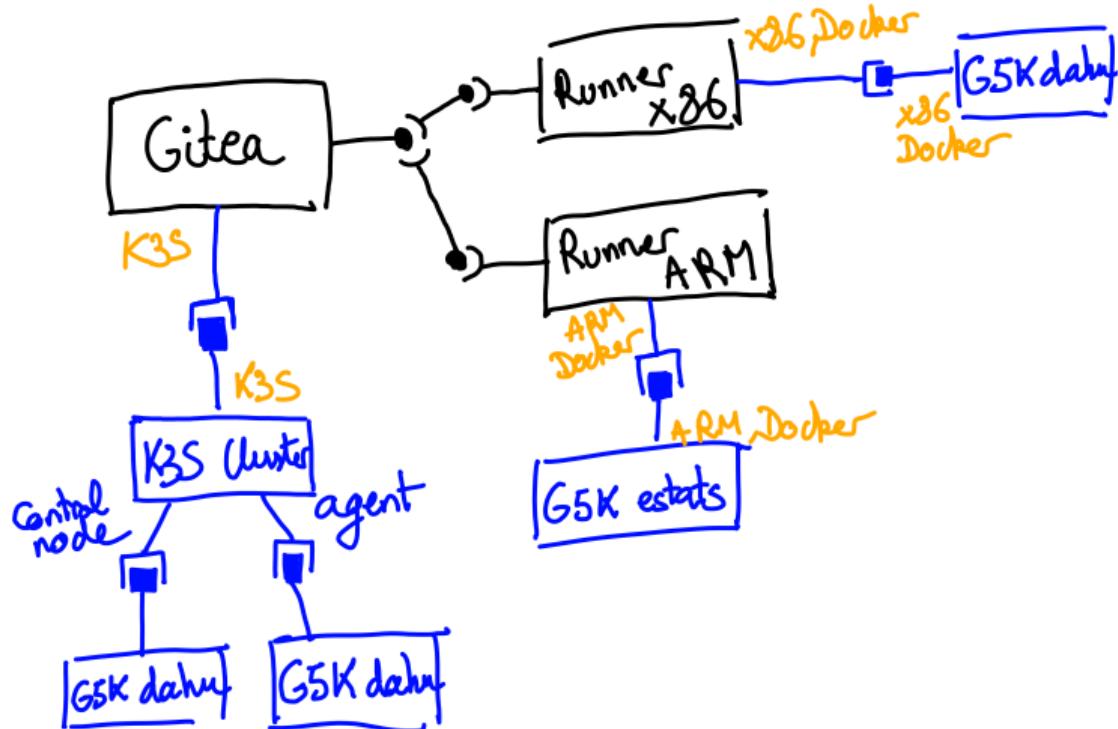
- Had to deploy own K3S cluster
- using one x86 node (dahu) for runner
- using one arm64 node (estats) for runner

## Scenario – Gitea and CI Runners – Production



- Had to deploy own K3S cluster
- using one x86 node (dahu) for runner
- using one arm64 node (estats) for runner

## Scenario – Gitea and CI Runners – Production



- Had to deploy own K3S cluster
- using one x86 node (dahu) for runner
- using one arm64 node (estats) for runner

## Scenario – Gitea and CI Runners – Production



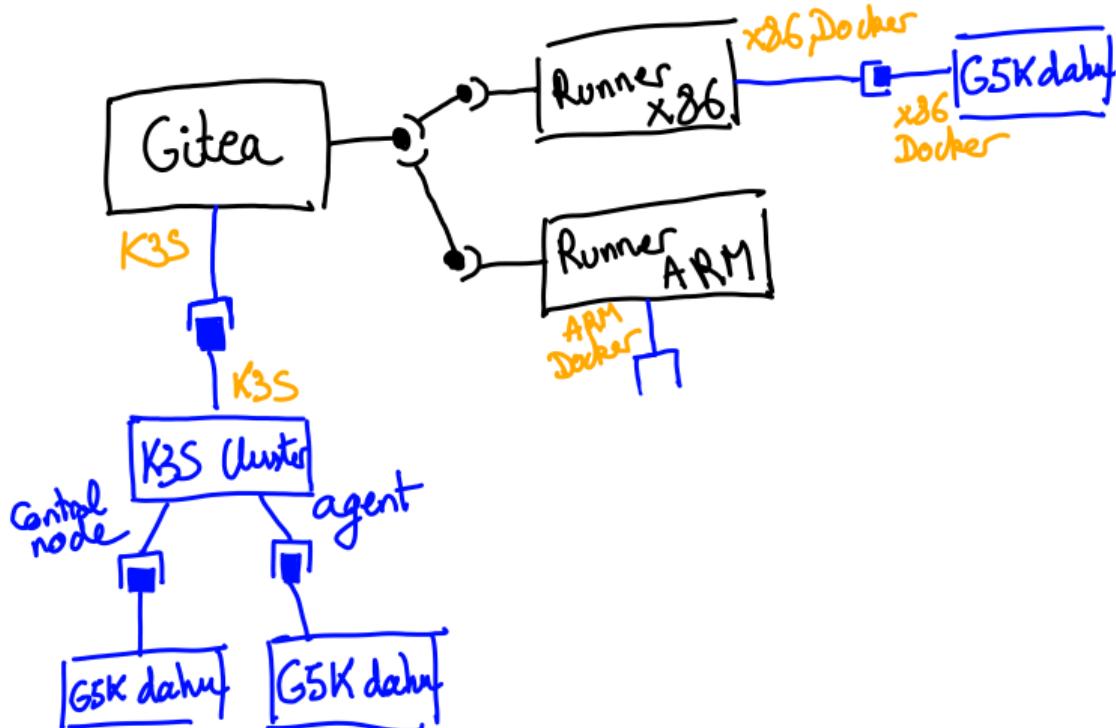
- Had to deploy own K3S cluster
- using one x86 node

Next Scenario: “Argh! No more arm64 node available...”



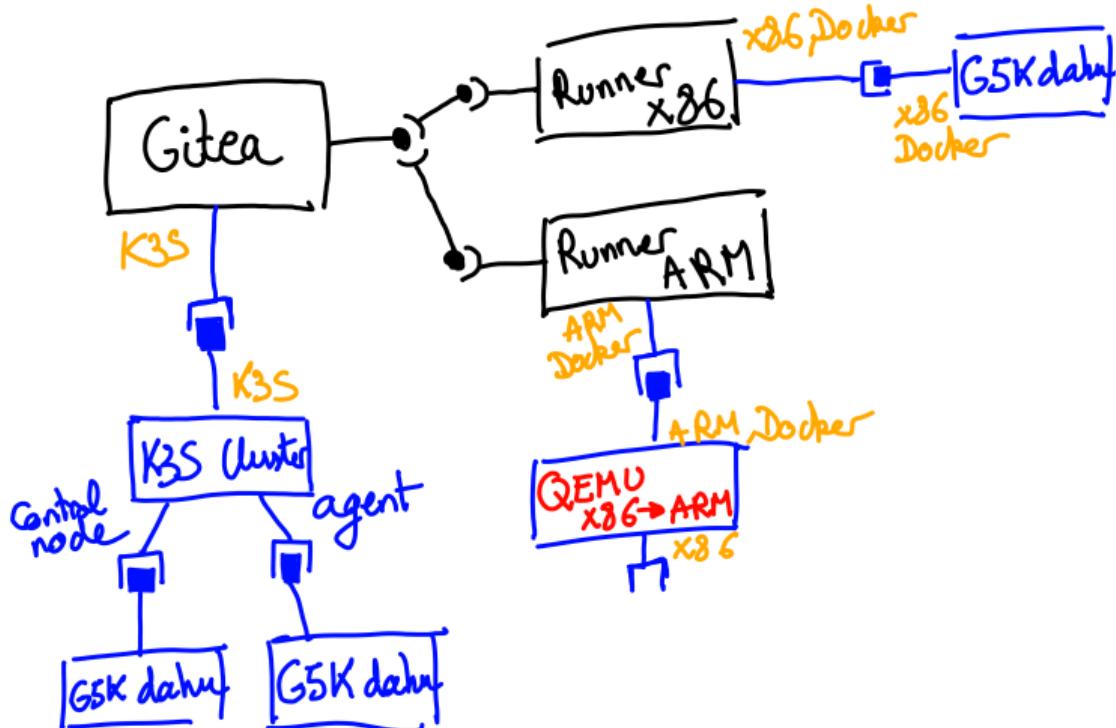
- using one arm64 node (estats) for runner

## Scenario – Gitea and CI Runners – Incomplete Production



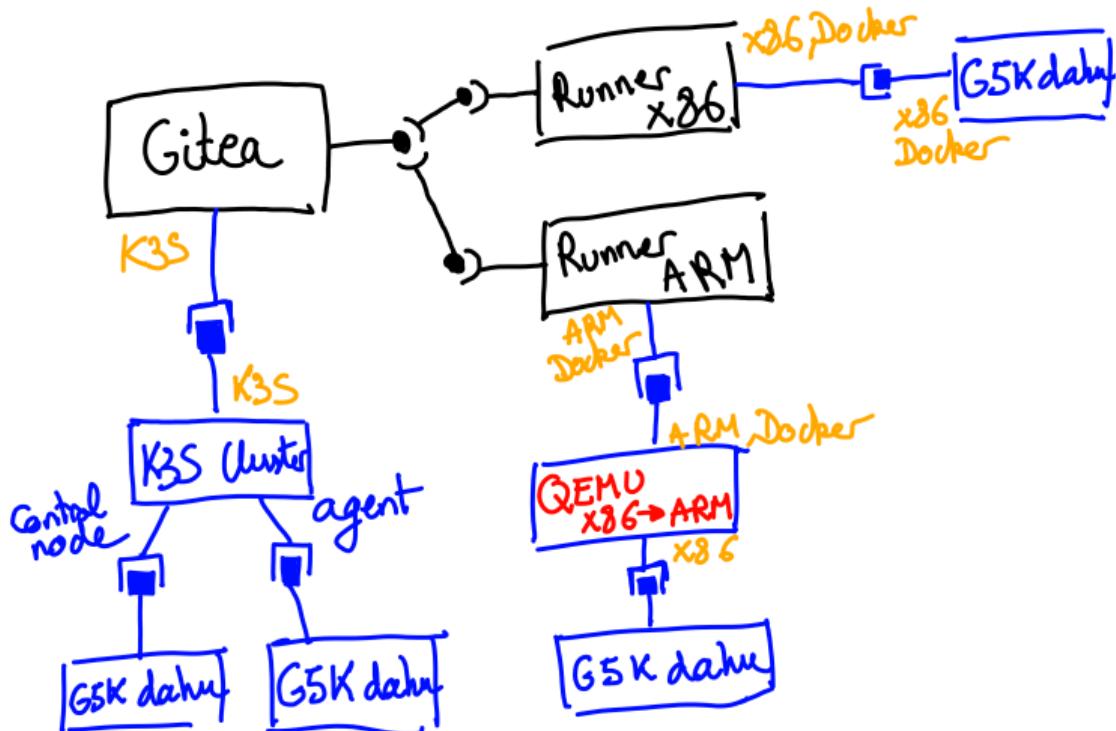
- no arm64 node available!
- use a x86 node + **QEMU**
- “deep encapsulation”!

## Scenario – Gitea and CI Runners – Incomplete Production



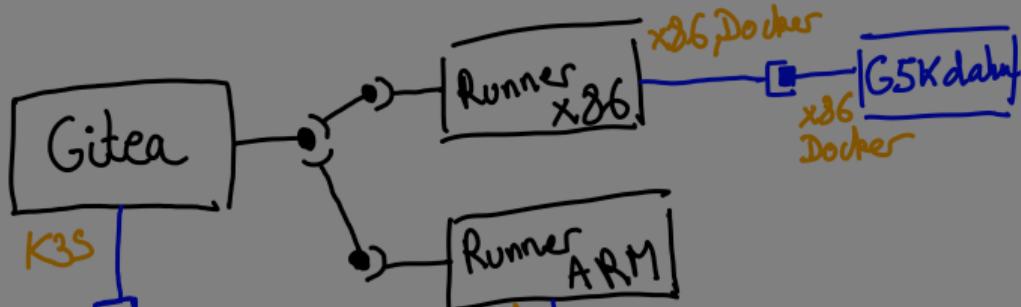
- no arm64 node available!
- use a x86 node + **QEMU**
- “deep encapsulation”!

## Scenario – Gitea and CI Runners – Incomplete Production



- no arm64 node available!
- use a x86 node + **QEMU**
- “deep encapsulation”!

## Scenario – Gitea and CI Runners – Incomplete Production



- no arm64 node available!

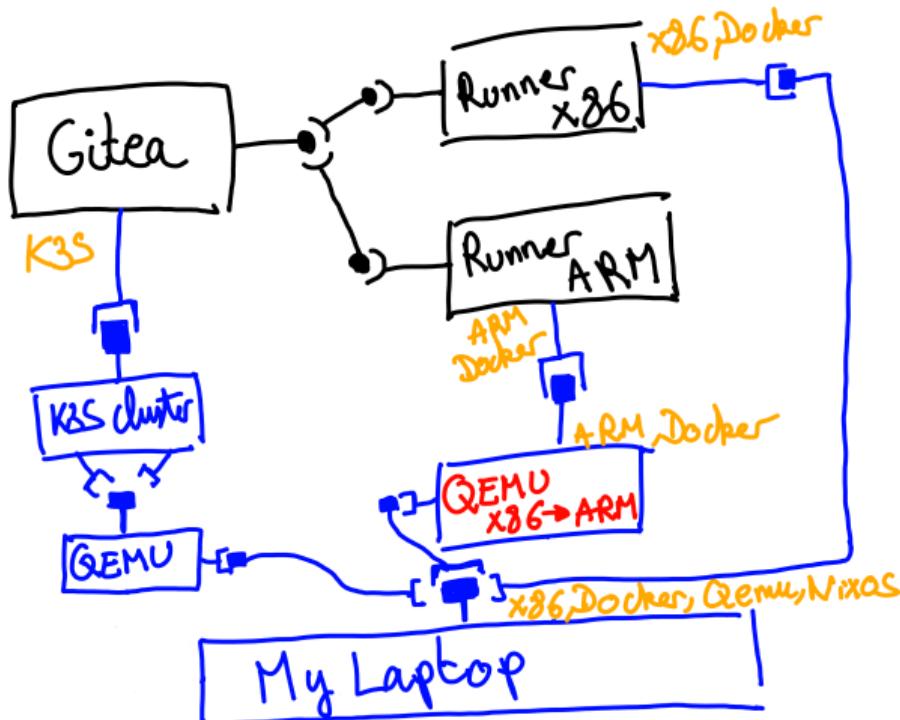
Next Scenario: “Let me test something on my laptop”



QEMU

- “deep encapsulation”!

## Scenario – Gitea and CI Runners – Development



- need to reduce the resource requirements (functional attributes) to fit on laptop
- K3S needs Debian
- VM for K3S cluster and VM for arm worker

## Conclusion & Perspectives

---

Pb: "How to model assemblies to support their execution in various contexts?"

## Conclusion and Analysis

- ✓ Model to express application assemblies independently from the resources
- ✓ Allow for portability of assemblies, and heterogeneous deployments
- ✓ Able to model complex deep "encapsulation levels"
- ✓ Can generate problem for a solver, and transform the output into Concerto
- ≈ Resource components not (yet) too difficult to implement in Concerto
- ≈ Might be difficult for components to communicate if "too deep" in encapsulation

# Conclusion & Perspectives

Pb: "How to model assemblies to support their execution in various contexts?"

## Conclusion and Analysis

It kinda works!

## Perspectives

- **Reconfiguration!**
- Support variants of components, composites, auto-wiring (?)
- Software environments
- Network, storage elements as resources ?
- Allow the Mapper to change the parameters of the resource components

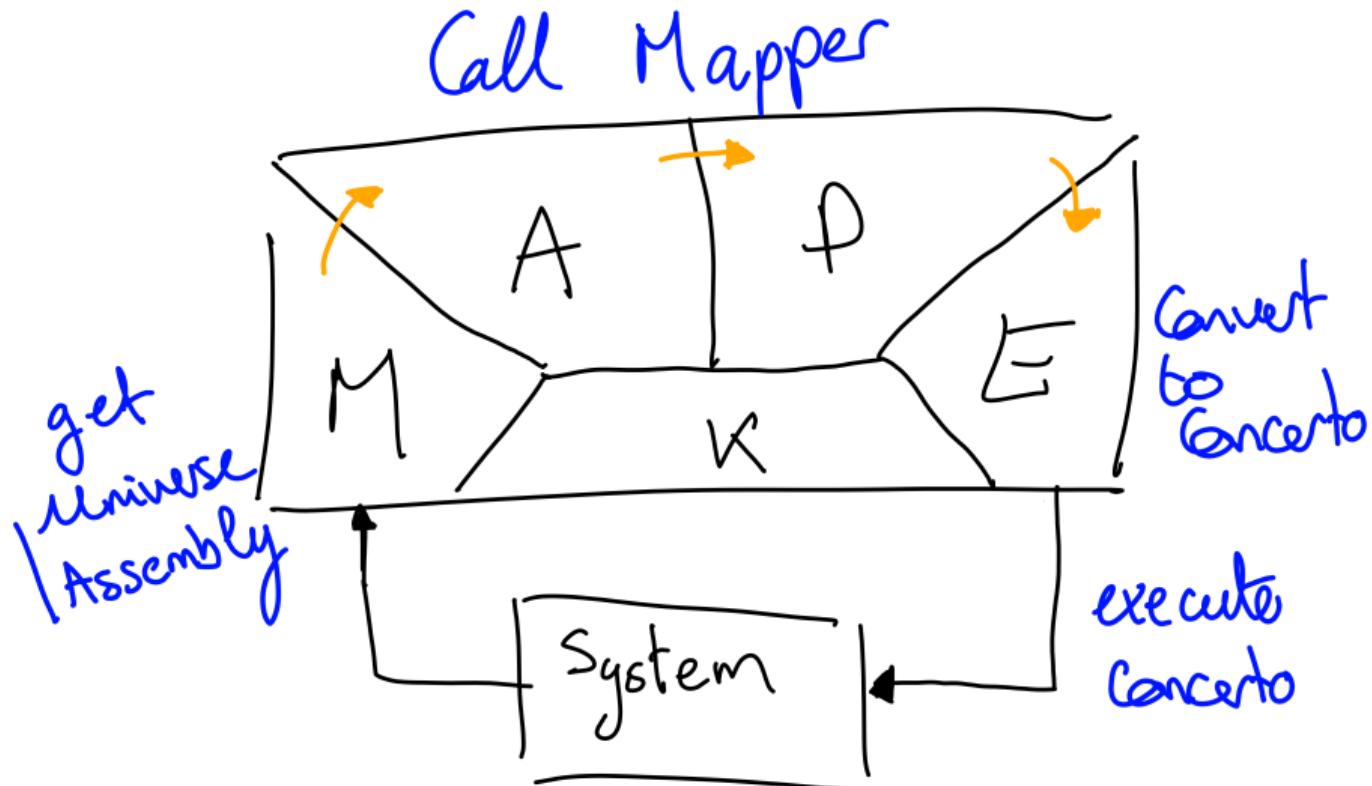
## Some Self-Reflection...

---

By making deployments easier/optimized, we also made **more** deployments possible...

→ Are we participating in the acceleration of the Cloud?

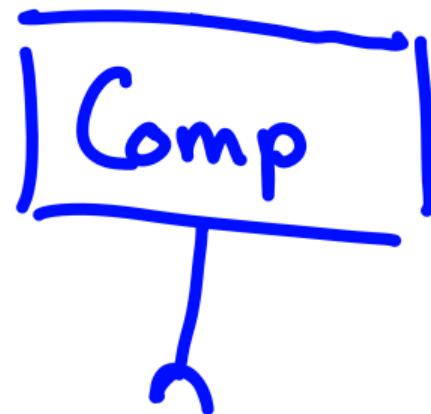
## The Global Picture



# Categories of Resource Components

## Applicative Components

- Uses one, or more, resources
- Does not provide any resource

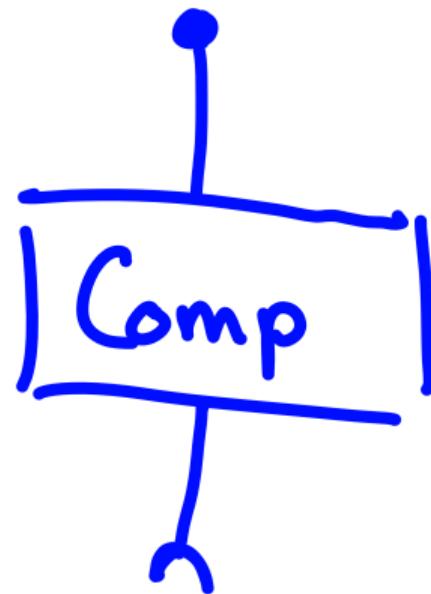


# Categories of Resource Components

## Applicative Components

## Transformers/Converters

- “Transform” one resource into another one
- Use one resource, provide one resource
- e.g., Virtual Machines, Containers,



# Categories of Resource Components

Applicative Components

Transformers/Converters

Resources Groups

- Gather several resources into a single interface
- e.g., Kubernetes



# Categories of Resource Components

Applicative Components

Transformers/Converters

Resources Groups

Allocations

- Already allocated resources
- Resources acquired via allocation
- Only provides resource(s)
- e.g., local machine, Grid'5000, AWS

