

Towards Reusability of Autonomic Controllers in High Performance Computing

GDR GPL - YODA, Vannes

Quentin GUILLOTEAU^{*}, Éric RUTTEN^{*}, Bogdan ROBU^{**}, Olivier RICHARD^{*}

^{*} Université Grenoble Alpes, Inria, CNRS, Grenoble INP, LIG

^{**} Université Grenoble Alpes, CNRS, Grenoble INP, GIPSA-lab

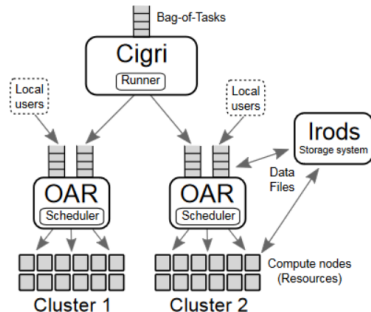
2022-06-08

Context

Idle HPC Resources \implies Lost Computing Power \rightsquigarrow **How to Harvest ?**

One Solution: *CiGri*

- **bag-of-tasks**: many, multi-parametric
- **Best-effort Jobs**: Lowest priority
- **Objective**: Collect grid idle resources



Problem

\nearrow Harvesting \implies \nearrow Perturbations (e.g. I/O) \rightsquigarrow **Trade-off**

\hookrightarrow Unpredictability \implies **runtime management**

Runtime management

Autonomic Computing and the MAPE-K Loop


Auto-regulating Systems given **high-level objectives**

Phases: **Monitor** \rightsquigarrow **Analyse** \rightsquigarrow **Plan** \rightsquigarrow **Execute** (with **K**nowledge)

Control Theory (Feedback Control Loop)

Regulate the behaviour of dynamical systems

\hookrightarrow Interpretation of the MAPE-K Loop



1 Introduction & Context

2 The *CiGri* Middleware

3 Design of Controllers

- Proportional-Integral
- Model-Free Control
- Ease of Design/Setup

4 Experimental Comparison

5 Conclusion & Perspectives

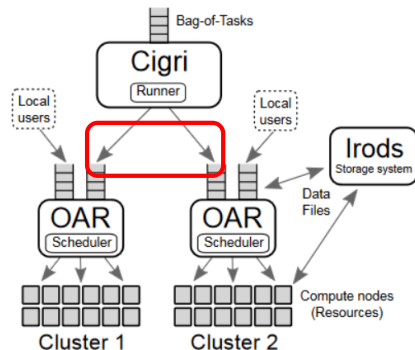
CiGri: Submission Loop (1/2)

Algorithm 1: Current Solution

```

rate = 3;
increase_factor = 1.5;
while tasks not executed in b-o-t do
  if no task running then
    submit rate tasks;
    rate = min(rate ×
      increase_factor, 100);
  end
  while nb of tasks running > 0
    do
      sleep during 30 sec;
    end
  end
end


```



CiGri: Submission (2/2)

The Issue

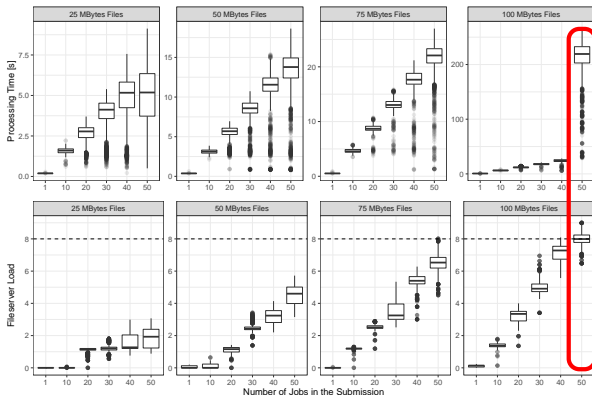
Must wait for termination of the previous submission to submit again
 ↪ reduce overload but introduce **underutilization** of the resources



Degradation of the File System Performances

↗ Jobs \Rightarrow ↗ I/O \Rightarrow ↗ More delay for users \rightsquigarrow **Perturbations**

Processing Time and Fileserver Load for different Submissions (number of jobs and filesize)



overload!

Sensor

- loadavg
- linear relation
- shows limits of FS
- estimation of perturbations

Our Global Problem and Objectives

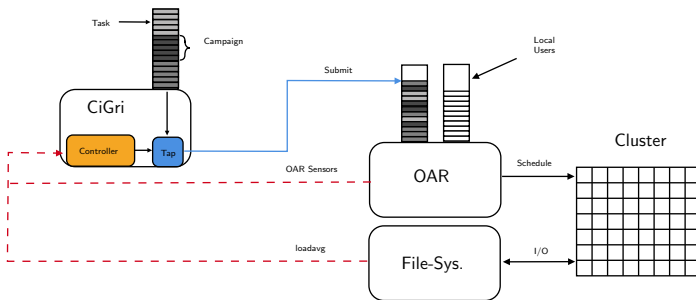
Objective

Harvest Idle Resources in a
non-intrusive way

- max cluster utilization
- min perturbations

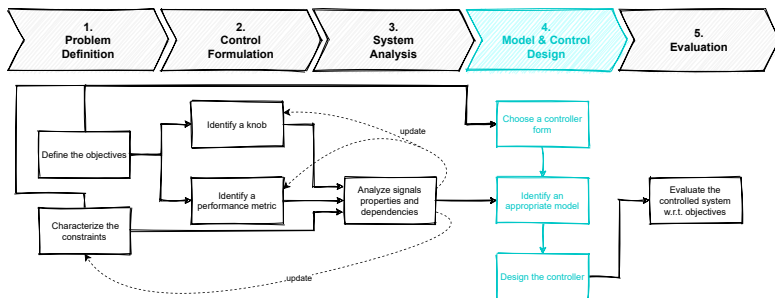
Means

- Instrumentation
 - **Actuator**: #jobs to submit, ...
 - **Sensor**: RJMS WQ, FS Load, ...
- **Controllers** (PID, RST, MFC, ...)
- Experimental Validation



Usual Method (e.g., PID) and Difficulties

↪ take into account current state of cluster \rightsquigarrow **use Control Theory**



However...

Cluster/Grid Administrators are **not** Control Theory experts

↪ **Design Cost? Setup Cost? Runtime Performances?**

Comparison Framework

Two Controllers

- Proportional-Integral (PI)
- Model-Free (MFC)

Variations: jobs (I/O, duration), cluster.


Reusability Criterion

- Design Time Cost
- Runtime Performances

Goal

Compare Controllers Reusability: Design Cost vs. Performances

QUALITATIVE COMPARISON OF DIFFERENT CONTROL SOLUTIONS



- 1 Introduction & Context
- 2 The *CiGri* Middleware
- 3 Design of Controllers
 - Proportional-Integral
 - Model-Free Control
 - Ease of Design/Setup
- 4 Experimental Comparison
- 5 Conclusion & Perspectives

PI: What are we looking for

First, a **Model** ... (i.e. how does the system behave (Open-Loop))

$$\mathbf{y}(k+1) = \sum_{i=0}^k a_i \mathbf{y}(k-i) + \sum_{j=0}^k b_j \mathbf{u}(k-j)$$

... then a **(PID) Controller** (i.e. the Closed-Loop behavior)

$$\text{Output} = \mathbf{K}_p \times \text{Error}_k + \mathbf{K}_i \times \sum_k \text{Error}_k + \mathbf{K}_d \times (\text{Error}_k - \text{Error}_{k-1})$$

Sensors & Actuators

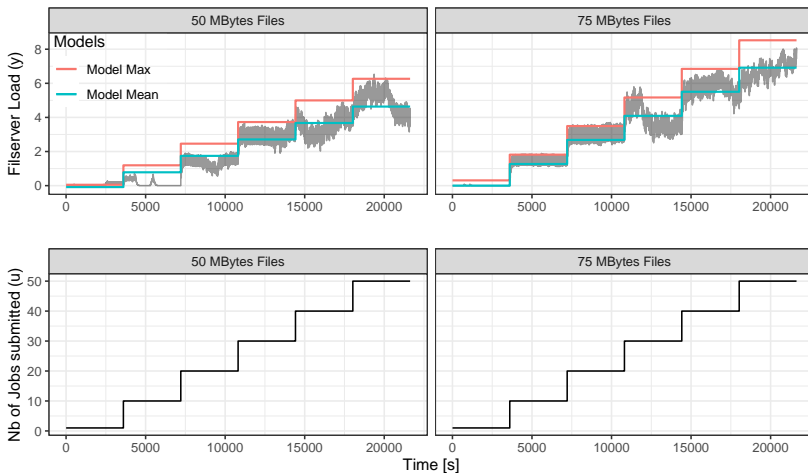
- Actuator: #jobs to sub $\rightsquigarrow \mathbf{u}$
- Sensor: FS Load $\rightsquigarrow \mathbf{y}$
- Error: *Reference* – *Sensor*

Method

- 1 Open-Loop expe (fixed \mathbf{u})
- 2 Model parameters (a_i, b_j)
- 3 Choice controller behavior (\mathbf{K}_*)

PI: Open-Loop and Identification

System Identification and (Linear) Model Fitting




$$y = \alpha + \beta_1 f + \beta_2 u + \gamma f u$$

PI: Closed-Loop Behavior

Open-Loop Experiments \Rightarrow Model (1st order) \Rightarrow Controller Gains K_p, K_i, K_d ,
 $y(k+1) = ay(k) + bu(k)$

Closed loop behaviour of our system for different values of (k_s, M_p)



Controller Gains are ...
 functions of the model and

- k_s : max **time** to steady state
- M_p : max **overshoot** allowed

Non-Intrusive Harvesting

- no overshoot
- but "fast" response

- 1 Introduction & Context
- 2 The *CiGri* Middleware
- 3 Design of Controllers
 - Proportional-Integral
 - **Model-Free Control**
 - Ease of Design/Setup
- 4 Experimental Comparison
- 5 Conclusion & Perspectives

What is Model-Free Control ? [Fliess & Join]

Model-Free Control

- Introduces *intelligent* Controllers (*iPID*)
- Easier to tune than PI
- Adapt to the plant/system
- can be equivalent to PI

$$\begin{cases} \hat{F}_k &= \frac{y_k - y_{k-1}}{\Delta t} - \alpha \times u_k \\ u_{k+1} &= \frac{-\hat{F}_k - \dot{y}_k^* + \mathbf{K}_p \times e_k}{\alpha} \end{cases}$$

- y_k : Load of File System
- u_k : #jobs *CiGri*
- \dot{y}_k^* : Derivative of ref. value

- \hat{F}_k : Estimation of the model
- α : **non-physical cst parameter**
- \mathbf{K}_p : **Gain of the controller**

Choice of Parameters

Empirical choice of parameters (α & \mathbf{K}_p)

- 1 Introduction & Context
- 2 The *CiGri* Middleware
- 3 Design of Controllers
 - Proportional-Integral
 - Model-Free Control
 - Ease of Design/Setup
- 4 Experimental Comparison
- 5 Conclusion & Perspectives

Ease of Design/Setup

Proportional-Integral

- Cumbersome to set up
 - Requires identification
 - Only for identified system
- + Behavior guarantee

Model-Free Control

- + Easy to set up
 - + No identification
 - + Should adapt to the plant
- No behavior guarantee

Take away

↔ MFC has lighter setup phase, but PI has more guarantees

1 Introduction & Context

2 The *CiGri* Middleware

3 Design of Controllers

- Proportional-Integral
- Model-Free Control
- Ease of Design/Setup

4 Experimental Comparison


5 Conclusion & Perspectives

Experimental Setup

Experimental Setup

- Experiments done on Grid'5000
- Emulation of a 100 node cluster
- 2 Intel Xeon E5-2630 v3
- CiGri jobs: sleep + write

Fileserver Load and Submission size

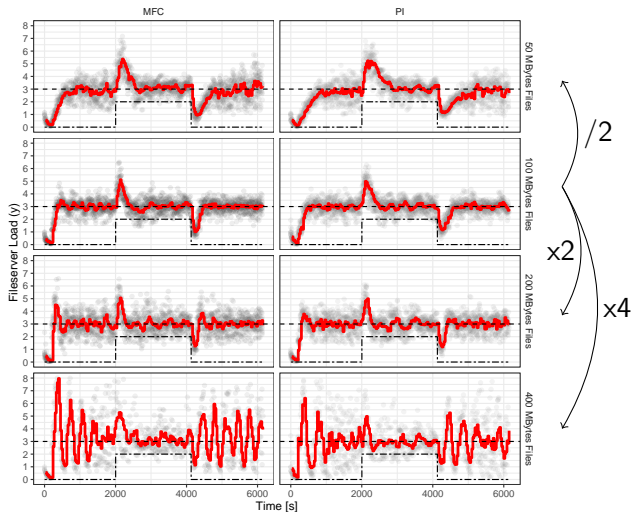


Synthetic Load

- Pure step
- Observe the ctrlr behavior:
 - response
 - oscillations

Variation in I/O

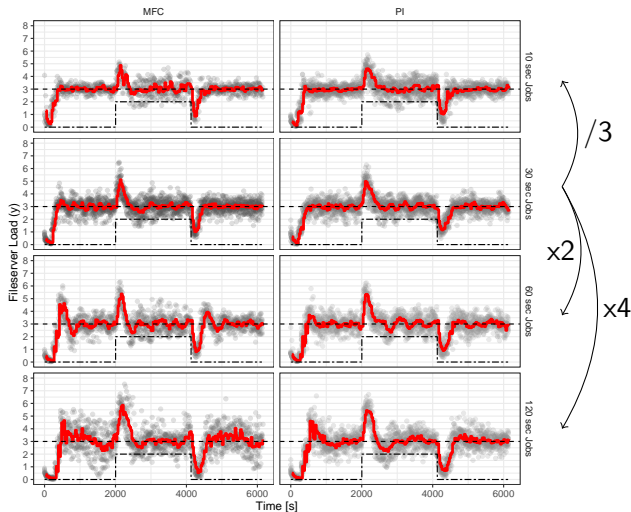
Comparison between the MFC and PI with variations in the I/O impact of jobs



- \simeq behavior
- MFC faster but more aggressive
- PI less variations for larger I/O

Variation in Execution Time

Comparison between the MFC and PI with variations in the Execution Times of jobs




- \simeq behavior
- MFC faster but more aggressive
- Job duration variations have less impact on control quality than the I/O quantity

Performances Comparison

Comparison of the Mean Control Errors for the Controllers with different Variations

99% confidence intervals



- 1 Introduction & Context
- 2 The *CiGri* Middleware
- 3 Design of Controllers
 - Proportional-Integral
 - Model-Free Control
 - Ease of Design/Setup
- 4 Experimental Comparison
- 5 Conclusion & Perspectives

Conclusion & Perspectives

Reminder of the Objective

Investigate the **Reusability** of Autonomic Controllers in HPC

Results

Compared 2 Controllers: PI & MFC on I/O and job dur. Variations

- MFC has smaller design cost
- \simeq performances for both controllers

\hookrightarrow **MFC seems more reusable than PI**

Perspectives

- Compare with other Solutions (e.g., PI Adaptive, MPC + GP)
- Investigate more variations dimensions (e.g., FS)